

**M A S A R Y K
U N I V E R S I T Y**

FACULTY OF INFORMATICS

Hardware-encrypted disks in Linux

Master's Thesis

ŠTĚPÁN HORÁČEK

Brno, Spring 2022

**M A S A R Y K
U N I V E R S I T Y**

FACULTY OF INFORMATICS

Hardware-encrypted disks in Linux

Master's Thesis

ŠTĚPÁN HORÁČEK

Advisor: Ing. Milan Brož, Ph.D.

Department of Computer Systems and Communications

Brno, Spring 2022



Declaration

Hereby I declare that this paper is my original authorial work, which I have worked out on my own. All sources, references, and literature used or excerpted during elaboration of this work are properly cited and listed in complete reference to the due source.

Štěpán Horáček

Advisor: Ing. Milan Brož, Ph.D.

Abstract

This is the abstract of my thesis, which can span multiple paragraphs.

Keywords

keyword1, keyword2, ...

Contents

1	Introduction	1
1.1	Thesis description	1
2	Hardware disk encryption	2
2.1	Self-encrypting drives	2
2.2	Inline encryption hardware	2
2.3	Software encryption	2
2.3.1	Linux stack	2
3	fscrypt	4
3.1	How to make it work?	4
3.2	How does it work?	5
3.2.1	Setup	5
3.2.2	Usage	6
4	TCG Opal 2.0	7
4.1	Technical stuff	7
4.1.1	Capability discovery	7
4.1.2	sessions	7
4.1.3	setup	8
4.1.4	set locking range	8
4.1.5	access...	8
4.1.6	8
4.2	Comparison with OPAL 1.0	8
4.3	Comparision with other TCG's SSC	8
5	Existing tools	9
5.1	Cryptsetup	9
5.2	sedutil	9
6	Security of hardware encryption	10
6.1	Attacks on hardware encryption	10
7	OPAL extension for cryptsetup	11

8 Conclusion	12
Bibliography	13

1 Introduction

Just keeping the citations here, like [1], [2], [3].

1.1 Thesis description

The thesis aims to analyze existing approaches to using hardware-encrypted block devices (disks) in Linux and propose integrating such devices into the cryptsetup project.

The implementation should use the Linux kernel interface.
Student should

- get familiar with and study available resources for self-encrypted drives, OPAL2 standard, block layer inline encryption,
- analyze and describe security of such drives,
- provide state-of-the-art overview of existing attacks,
- implement proof-of-concept extension to cryptsetup project (and possibly propose changes for Linux kernel).

The student should be familiar with C code for low-level system programming and cryptography concepts.

2 Hardware disk encryption

Hardware disk encryption is a technology that provides confidentiality of data stored on a storage device using encryption provided by hardware.

Some general overview about what it is, I guess., maybe talk about both sw and hw enc instead

Describe generic stuff: provisioning, locking, key types DEK, KEK, MEK, processes, locking ranges vs. FDE, ... keyslots?,,, (and for opal TPer, SP, ..)

mention LUKS somewhere...

The disk encryption process can be conducted in logically and physically different places, depending on the type. In the following sections of this chapter, we will further describe three such types.

2.1 Self-encrypting drives

at least try to mention ATA security,,, which means also mentioning opal,,, and TCG Enterprise...

2.2 Inline encryption hardware

Compared to the previously described self-encrypting drives, inline encryption hardware is separated from the disk, and

something about how it provides actually a way to check the encrypted content, right?

2.3 Software encryption

probably just a quick overview, ... maybe change the chapter to just "disk encryption"... or mention this just shortly at the start...

2.3.1 Linux stack

rethink where to put this... maybe tools?

dm-crypt

fscrypt

3 fscrypt

Assuming ext4 in this chapter...

3.1 How to make it work?

As of now, blk-layer inline encryption is supported only by two filesystems in Linux: ext4 and F2FS. Need to use a kernel with `CONFIG_FS_ENCRYPTION_INLINE_CRYPT` enabled. Need to first mount the filesystem with the inline encryption flag:

```
mount -t ext4 /dev/foo /mnt/foo -o inlinecrypt
```

It is not enough to specify the inline encryption flag, the encryption itself also must be enabled. Assuming ext4 filesystem, the encryption can be enabled after mounting like so:

```
tune2fs -O encrypt "/dev/loop0"
```

After this, fscrypt can be used as normal and it will use inline encryption for this filesystem.

In order to encrypt a folder using a fscrypt, the following must be done: an encryption key must be added and the encryption policy must be created.

```
int fd = open(pathname, O_RDONLY | O_CLOEXEC);

struct fscrypt_add_key_arg *key_request = calloc
    (1, sizeof(struct fscrypt_add_key_arg) +
     key_len);
struct fscrypt_policy_v2 policy_request = { 0 };

// add a key
key_request->key_spec.type =
    FSCRYPT_KEY_SPEC_TYPE_IDENTIFIER;
key_request->key_id = 0;
key_request->raw_size = key_len;
memcpy(key_request->raw, key, key_len);
ioctl(fd, FS_IOC_ADD_ENCRYPTION_KEY, key_request
    );
```

```
// set a policy
policy_request.version = 2;
policy_request.contents_encryption_mode =
    FSCRYPT_MODE_AES_256_XTS;
policy_request.filename_encryption_mode =
    FSCRYPT_MODE_AES_256_CTS;
policy_request.flags =
    FSCRYPT_POLICY_FLAGS_PAD_8 |
    FSCRYPT_POLICY_FLAGS_PAD_16 |
    FSCRYPT_POLICY_FLAGS_PAD_32;
memcpy(policy_request.master_key_identifier,
    key_request->key_spec.u.identifier,
    FSCRYPT_KEY_IDENTIFIER_SIZE);
ioctl(fd, FS_IOC_SET_ENCRYPTION_POLICY, &
    policy_request);
```

This code sets up an encryption policy for the file specified by the pathname.

3.2 How does it work?

3.2.1 Setup

During mounting the “inlinecrypt”/SB_INLINECRYPT flag is written into the super_block structure.

It all starts in `__ext4_new_inode`. This is the internal function used when creating new inodes, called by functions such as `ext4_create` when creating a new file.

The function (if it is not inode used for large extended attributes?) calls `fscrypt_prepare_new_inode`.

```
fscrypt_prepare_new_inode -> fscrypt_setup_encryption_info ->
setup_file_encryption_key -> fscrypt_select_encryption_impl
```

In `fscrypt_select_encryption_impl` there is actually the only place where the SB_INLINECRYPT flag is used. ... Calls `blk_crypto_config_supported` to check the device’s crypto profile. Afterwards, `fscrypt_select_encryption_impl` function sets the `(fscrypt_info *)ci->ci_inlinecrypt`.

`setup_per_mode_enc_key` then sets the `(fscrypt_info *)ci->ci_enc_key`.

3.2.2 Usage

The bio function is stored in

```
(struct bio *)bio->(struct bio_crypt_ctx *)bi_crypt_context
```

Function `fscrypt_set_bio_crypt_ctx` changes the file's bio to use inline encryption... simply calls the blk layer `bio_crypt_set_ctx`.

Calling `submit_bio` like normally ... `__submit_bio` calls `__blk_crypto_bio_prep`...

"If the bio crypt context provided for the bio is supported by the underlying device's inline encryption hardware, do nothing."

`__blk_crypto_rq_bio_prep` however sets the context of the request to the one of the bio... After the bio prep `blk_mq_submit_bio` gets called (which calls `blk_mq_bio_to_request`, and after that also `blk_crypto_init_request->blk_crypto_get_keyslot` which updates the devices keyslot to contain the new key..., but does nothing if the device does not have keyslots)..... the info about the key to use then has to be acquired by the driver from request->

Most important are probably structures `blk_crypto_ll_ops` and `blk_crypto_profile`... just two operations, program key and evict key.

how to get crypto profile from outside..

Where does the hardware come to play?

`ufshcd_exec_raw_upiu_cmd()->ufshcd_issue_devman_upiu_cmd()->ufshcd_prepare` sets the header with the correct keyslot.

4 TCG Opal 2.0

TCG Opal SSC (Security Subsystem Class) 2.0 is a specification for storage devices, aiming to provide confidentiality of stored data while the disk implementing this specification is turned off [1]. The Opal standard is one of the representants of the self-encrypting drive approach to hardware disk encryption.

developed by TCG (Trusted Computing Group)

maybe move under the SED chapter, but it is probably going to be too big for that..., and important

structure of the standards described in https://trustedcomputinggroup.org/wp-content/uploads/TCGandNVMe_Joint_White_Paper-TCG_Storage_Opal_and_NVMe_FINAL.pdf.

4.1 Technical stuff

I expect stuff here like important headers, codes, ...

SP = service provider, tables and methods that operate upon them

4.1.1 Capability discovery

begin with Level 0 discovery: (4, p. 3.3.6) specifies the general security receives command for all TCG storage devices [1, p. 3.1.1] specifies the Opal 2 specific parts of the header.

4.1.2 sessions

In order to facilitate communication between the host and SP, sessions must be used. During these sessions, methods are used...

example of method:

```
SMUID.StartSession [
HostSessionID : uinteger,
SPID : uidref {SPObjectUID},
Write : boolean,
HostChallenge = bytes,
HostExchangeAuthority = uidref {
    AuthorityObjectUID},
```

```
HostExchangeCert = bytes ,
HostSigningAuthority = uidref {
    AuthorityObjectUID},
HostSigningCert = bytes ,
SessionTimeout = uinteger ,
TransTimeout = uinteger ,
InitialCredit = uinteger ,
SignedHash = bytes ]
=>
SMUID.SyncSession [ see SyncSession definition
    in 5.2.3.2]
```

Each method call is defined by the caller UID, method UID, and values of its parameters. The method calls are coded using tokens... start list, end list, optional argument start, ... short int, long int,....

To call a method use IF-SEND command, to get result of the method call use IF-RECV command.

4.1.3 setup

4.1.4 set locking range

4.1.5 access...

4.1.6 ...

4.2 Comparison with OPAL 1.0

OPAL 2 disks are *optionally* backwards compatible.

cite https://trustedcomputinggroup.org/wp-content/uploads/TCG_Storage-Opal_SSC_FAQ.pdf

4.3 Comparision with other TCG's SSC

Enterprise, Opalite, Pyrite, Ruby

5 Existing tools

5.1 Cryptsetup

... tool for disk encryption setup.
nothing for opal, right?

5.2 sedutil

supports opal, but the code seems to be abandoned

6 Security of hardware encryption

specify threat model, ,,,,,, offline, online,,, offline multiple times???
.. maybe change position of this chapter, since it's splitting tools
and change to a tool

6.1 Attacks on hardware encryption

7 OPAL extension for cryptsetup

8 Conclusion

Bibliography

1. *TCG Storage Security Subsystem Class: Opal*. 2022-01. Version 2.02. Standard. Trusted Computing Group. Available also from: <https://trustedcomputinggroup.org/resource/storage-work-group-storage-security-subsystem-class-opal>.
2. BIGGERS, Eric; TANGIRALA, Satya. *Inline Encryption*. 2021. Available also from: <https://kernel.org/doc/html/v5.17/block/inline-encryption.html>.
3. MÜLLER, Tilo; LATZO, Tobias; FREILING, Felix C. *Self-Encrypting Disks pose Self-Decrypting Risks: How to break Hardware-based Full Disk Encryption*. Available also from: <https://fau1-files.cs.fau.de/filepool/projects/sed/seds-at-risks.pdf>.
4. *TCG Storage Architecture Core Specification*. Tech. rep. Available also from: <https://trustedcomputinggroup.org/resource/tcg-storage-architecture-core-specification>.