

**M A S A R Y K
U N I V E R S I T Y**

FACULTY OF INFORMATICS

Hardware-encrypted disks in Linux

Master's Thesis

ŠTĚPÁN HORÁČEK

Brno, Spring 2022

**M A S A R Y K
U N I V E R S I T Y**

FACULTY OF INFORMATICS

Hardware-encrypted disks in Linux

Master's Thesis

ŠTĚPÁN HORÁČEK

Advisor: Ing. Milan Brož, Ph.D.

Department of Computer Systems and Communications

Brno, Spring 2022



Declaration

Hereby I declare that this paper is my original authorial work, which I have worked out on my own. All sources, references, and literature used or excerpted during elaboration of this work are properly cited and listed in complete reference to the due source.

Štěpán Horáček

Advisor: Ing. Milan Brož, Ph.D.

Abstract

This is the abstract of my thesis, which can span multiple paragraphs.

Keywords

keyword1, keyword2, ...

Contents

1	Introduction	1
1.1	Thesis description	1
2	Hardware disk encryption	2
2.1	Self-encrypting drives	2
2.2	Inline encryption hardware	2
2.3	Software encryption	2
2.3.1	Linux stack	2
3	fscrypt	4
3.1	How to make it work?	4
3.2	How does it work?	5
3.2.1	Setup	5
3.2.2	Usage	6
4	TCG Opal 2.0	7
4.1	Technical stuff	7
4.1.1	Capability discovery	7
4.1.2	sessions	7
4.1.3	setup	8
4.1.4	set locking range	8
4.1.5	access...	8
4.1.6	9
4.2	Comparison with OPAL 1.0	9
4.3	Comparison with other TCG's SSC	9
5	Existing tools	10
5.1	Cryptsetup	10
5.2	sedutil	10
6	Security of hardware encryption	11
6.1	Attacks on hardware encryption	11
7	OPAL extension for cryptsetup	12

8	Conclusion	13
9	TCG Opal 2.0	14
9.1	Structure of the standard	14
9.2	Architecture	15
9.2.1	Trusted Peripheral	15
9.2.2	Security Provider	15
9.3	Capability discovery	15
9.3.1	Level 0 discovery	15
9.3.2	Level 1 discovery	16
9.3.3	Level 2 discovery	17
9.4	Life cycle	18
9.4.1	SP Issuance	18
9.4.2	Life cycles in Opal	18
9.5	Communication	19
9.5.1	Method calls	19
9.5.2	Sessions	20
9.5.3	Transactions	21
9.5.4	Optional features	21
9.6	Features	22
9.6.1	Single User Mode	22
9.6.2	MBR shadowing	22
9.6.3	Access Control	22
9.7	Templates	23
9.7.1	Base template	23
9.7.2	Admin template	23
9.7.3	Crypto template	24
9.7.4	Locking template	24
9.7.5	Clock template	24
9.7.6	Log template	24
9.8	Admin SP	24
9.8.1	Methods	24
9.9	Locking SP	25
9.9.1	Locking range	25
9.10	Implementations/Direct usage/something like this . . .	25
9.10.1	Linux ioctl	25
9.10.2	Direct communication	27

10 Notes	28
10.0.1 Templates	32
11 Data	33
11.0.1 Discovery	33
Bibliography	35

1 Introduction

Just keeping the citations here, like [1], [2], [3].

1.1 Thesis description

The thesis aims to analyze existing approaches to using hardware-encrypted block devices (disks) in Linux and propose integrating such devices into the cryptsetup project.

The implementation should use the Linux kernel interface.
Student should

- get familiar with and study available resources for self-encrypted drives, OPAL2 standard, block layer inline encryption,
- analyze and describe security of such drives,
- provide state-of-the-art overview of existing attacks,
- implement proof-of-concept extension to cryptsetup project (and possibly propose changes for Linux kernel).

The student should be familiar with C code for low-level system programming and cryptography concepts.

2 Hardware disk encryption

Hardware disk encryption is a technology that provides confidentiality of data stored on a storage device using encryption provided by hardware.

Some general overview about what it is, I guess., maybe talk about both sw and hw enc instead

Describe generic stuff: provisioning, locking, key types DEK, KEK, MEK, processes, locking ranges vs. FDE, ... keyslots?,,, (and for opal TPer, SP, ..)

mention LUKS somewhere...

The disk encryption process can be conducted in logically and physically different places, depending on the type. In the following sections of this chapter, we will further describe three such types.

2.1 Self-encrypting drives

at least try to mention ATA security,,, which means also mentioning opal,,, and TCG Enterprise...

2.2 Inline encryption hardware

Compared to the previously described self-encrypting drives, inline encryption hardware is separated from the disk, and

something about how it provides actually a way to check the encrypted content, right?

2.3 Software encryption

probably just a quick overview, ... maybe change the chapter to just "disk encryption"... or mention this just shortly at the start...

2.3.1 Linux stack

rethink where to put this... maybe tools?

dm-crypt

fscrypt

3 fscrypt

Assuming ext4 in this chapter...

3.1 How to make it work?

As of now, blk-layer inline encryption is supported only by two filesystems in Linux: ext4 and F2FS. Need to use a kernel with CONFIG_FS_ENCRYPTION_INLINE_CRYPT enabled. Need to first mount the filesystem with the inline encryption flag:

```
mount -t ext4 /dev/foo /mnt/foo -o inlinecrypt
```

It is not enough to specify the inline encryption flag, the encryption itself also must be enabled. Assuming ext4 filesystem, the encryption can be enabled after mounting like so:

```
tune2fs -O encrypt "/dev/loop0"
```

After this, fscrypt can be used as normal and it will use inline encryption for this filesystem.

In order to encrypt a folder using a fscrypt, the following must be done: an encryption key must be added and the encryption policy must be created.

```
int fd = open(pathname, O_RDONLY | O_CLOEXEC);

struct fscrypt_add_key_arg *key_request = calloc
    (1, sizeof(struct fscrypt_add_key_arg) +
    key_len);
struct fscrypt_policy_v2 policy_request = { 0 };

// add a key
key_request->key_spec.type =
    FSCRYPT_KEY_SPEC_TYPE_IDENTIFIER;
key_request->key_id = 0;
key_request->raw_size = key_len;
memcpy(key_request->raw, key, key_len);
ioctl(fd, FS_IOC_ADD_ENCRYPTION_KEY, key_request
    );
```

```
// set a policy
policy_request.version = 2;
policy_request.contents_encryption_mode =
    FSCRYPT_MODE_AES_256_XTS;
policy_request.filename_encryption_mode =
    FSCRYPT_MODE_AES_256_CTS;
policy_request.flags =
    FSCRYPT_POLICY_FLAGS_PAD_8 |
    FSCRYPT_POLICY_FLAGS_PAD_16 |
    FSCRYPT_POLICY_FLAGS_PAD_32;
memcpy(policy_request.master_key_identifier,
    key_request->key_spec.u.identifier,
    FSCRYPT_KEY_IDENTIFIER_SIZE);
ioctl(fd, FS_IOC_SET_ENCRYPTION_POLICY, &
    policy_request);
```

This code sets up an encryption policy for the file specified by the pathname.

3.2 How does it work?

3.2.1 Setup

During mounting the “inlinecrypt”/SB_INLINECRYPT flag is written into the super_block structure.

It all starts in `__ext4_new_inode`. This is the internal function used when creating new inodes, called by functions such as `ext4_create` when creating a new file.

The function (if it is not inode used for large extended attributes?) calls `fscrypt_prepare_new_inode`.

```
fscrypt_prepare_new_inode -> fscrypt_setup_encryption_info ->
setup_file_encryption_key -> fscrypt_select_encryption_impl
```

In `fscrypt_select_encryption_impl` there is actually the only place where the SB_INLINECRYPT flag is used. ... Calls `blk_crypto_config_supported` to check the device’s crypto profile. Afterwards, `fscrypt_select_encryption_impl` function sets the `(fscrypt_info *)ci->ci_inlinecrypt`.

`setup_per_mode_enc_key` then sets the `(fscrypt_info *)ci->ci_enc_key`.

3.2.2 Usage

The bio function is stored in

`(struct bio *)bio->(struct bio_crypt_ctx *)bi_crypt_context`

Function `fscrypt_set_bio_crypt_ctx` changes the file's bio to use inline encryption... simply calls the blk layer `bio_crypt_set_ctx`.

Calling `submit_bio` like normally ... `__submit_bio` calls `__blk_crypto_bio_prep...`

"If the bio crypt context provided for the bio is supported by the underlying device's inline encryption hardware, do nothing."

`__blk_crypto_rq_bio_prep` however sets the context of the request to the one of the bio... After the bio prep `blk_mq_submit_bio` gets called (which calls `blk_mq_bio_to_request`, and after that also `blk_crypto_init_request->blk_crypto_get_keyslot` which updates the devices keyslot to contain the new key..., but does nothing if the device does not have keyslots)..... the info about the key to use then has to be acquired by the driver from request->

Most important are probably structures `blk_crypto_ll_ops` and `blk_crypto_profile`... just two operations, program key and evict key.

how to get crypto profile from outside..

Where does the hardware come to play?

`ufshcd_exec_raw_upiu_cmd()->ufshcd_issue_devman_upiu_cmd()->ufshcd_prepare` sets the header with the correct keyslot.

4 TCG Opal 2.0

TCG Opal SSC (Security Subsystem Class) 2.0 is a specification for storage devices, aiming to provide confidentiality of stored data while the disk implementing this specification is turned off [1]. The Opal standard is one of the representants of the self-encrypting drive approach to hardware disk encryption.

developed by TCG (Trusted Computing Group)

maybe move under the SED chapter, but it is probably going to be too big for that..., and important

structure of the standards described in https://trustedcomputinggroup.org/wp-content/uploads/TCGandNVMe_Joint_White_Paper-TCG_Storage_Opal_and_NVMe_FINAL.pdf.

4.1 Technical stuff

I expect stuff here like important headers, codes, ...

SP = service provider, tables and methods that operate upon them

4.1.1 Capability discovery

begin with Level 0 discovery: (4, p. 3.3.6) specifies the general security receives command for all TCG storage devices [1, p. 3.1.1] specifies the Opal 2 specific parts of the header.

The Opal 2 feature header contains several important information usable in the future usage of the device. Most importantly the base ComID [4, p. 3.3.2]. The base ComID, together with another parameter the number of ComIDs, define the range of possible static ComIDs.

ComID are divided into static and dynamic (managed by ComID management).

4.1.2 sessions

In order to facilitate communication between the host and SP, sessions must be used. During these sessions, methods are used...

example of method:

`SMUID.StartSession [`

```

HostSessionID : uinteger,
SPID : uidref {SPObjectUID},
Write : boolean,
HostChallenge = bytes,
HostExchangeAuthority = uidref {
    AuthorityObjectUID},
HostExchangeCert = bytes,
HostSigningAuthority = uidref {
    AuthorityObjectUID},
HostSigningCert = bytes,
SessionTimeout = uinteger,
TransTimeout = uinteger,
InitialCredit = uinteger,
SignedHash = bytes ]
=>
SMUID.SyncSession [ see SyncSession definition
    in 5.2.3.2]

```

Each method call is defined by the caller UID, method UID, and values of its parameters. The method calls are coded using tokens... start list, end list, optional argument start, ... short int, long int,....

To call a method use IF-SEND command, to get result of the method call use IF-RECV command.

In order to authenitize the user, one can use many different ...

Methods are sent using packets. There are three types of packets: ComPackets, packets and subpackets. A single ComPacket can contain multiple packets and single packet can contain multiple subpackets. Each ComPacket is associated with a single ComID, each packet is associated with a single session.

4.1.3 setup

4.1.4 set locking range

global locking range is a special locking range that protects LBAs not protected by any of the other locking ranges.

4.1.5 access...

4.1.6 ...

4.2 Comparison with OPAL 1.0

OPAL 2 disks are *optionally* backwards compatible.

cite https://trustedcomputinggroup.org/wp-content/uploads/TCG_Storage-Opal_SSC_FAQ.pdf

4.3 Comparision with other TCG's SSC

Enterprise, Opalite, Pyrite, Ruby

5 Existing tools

5.1 Cryptsetup

... tool for disk encryption setup.
nothing for opal, right?

5.2 sedutil

supports opal, but the code seems to be abandoned

6 Security of hardware encryption

specify threat model, ,,,,,, offline, online,,, offline multiple times???
.. maybe change position of this chapter, since it's splitting tools
and change to a tool

6.1 Attacks on hardware encryption

7 OPAL extension for cryptsetup

8 Conclusion

9 TCG Opal 2.0

TCG Opal Security Subsystem Class (SSC) 2.0 (hereinafter referred to simply as “Opal”) is a specification for storage devices, aiming to provide confidentiality of stored data while the conforming disk is turned off [1]. It is one of the representatives of the self-encrypting drive approach to hardware disk encryption.

Developed in 2012 by Trusted Computing Group (TCG) as a successor of Opal 1.0.

Information for this chapter comes primarily from the Opal standards, that is the Core standard [4] and the Opal standard [1].

9.1 Structure of the standard

The Opal 2.0 standard is defined as a subsystem extending the TCG Storage Architecture Core Specification standard. The TCG Storage Architecture Core Specification standard [4] specifies the core features and properties shared among several different types of storage security subsystems, that extend the core functionality by specifying additional features or define the set of mandatory features. Each of these subsystems is focused on a different use case. These subsystems are namely:

- Opal — targeted at a corporate and personal usage. Described more closely in the rest of this chapter.
- Opalite — simplified Opal. Does not mandate features such as locking ranges, decreases the minimal number of admin and user authorities [5].
- Pyrite — encryption-less Opalite. Similar to Opalite, however it does not mandate encryption, and instead may offer only logical access control [6].
- Ruby — focused on data centers and server drives. Offers only global range encryption, weaker configuration of access control, no pre-boot authentication support [7]. Replacing the older Enterprise subsystem.

9.2 Architecture

The standard defines several parts of the trusted device.

9.2.1 Trusted Peripheral

Trusted peripheral (TPer) is a device located on the disk that provides the security of the data on disk. A TPer consists of multiple Security Providers.

9.2.2 Security Provider

Security Provider (SP) is defined as a set of tables, methods and an access control. Each SP is derived from a set of templates. These templates define a set of the tables and methods, aimed at one functionality, subset of which the SP implements. The templates are described more closely in later chapter ???. The Opal 2.0 standard defines that at least the Admin SP and Locking SP must be present in the TPer. The Admin SP tasked with administrating the TPer and other SPs, which may include creating new SPs, deleting existing SPs, or providing information about SPs. The Locking SP provides access to functionality such as managing locking ranges, locking the drive, or managing access control. Both of these SPs are described more closely in later chapter ??.

9.3 Capability discovery

In order to find out the properties and features of a particular device, there exists the so called Discovery process. This process is divided into three levels, each with different reported information and a different approach to access the information.

9.3.1 Level 0 discovery

Level 0 discovery provides basic information about the secure device, and is performed using only the secure send and secure receive commands of the device. The information is provided through feature descriptor. The presence of the feature descriptor header means that

the feature is supported and the fields of the header describe the basic properties of that feature.

Some of the features described by these descriptors are the TPer feature (supported communication features such as ack/nack support, comid management, buffer management, async communication, ...), Locking (whether disk is locked, etc.), Geometry feature (parameters of the disk, such as block size), Opal V1.0 feature, SingleUser feature (...), DataStore (size of the table), OPAL 2.0 feature (base ComID, number of ComIDs, default pin, number of locking users/admins).

9.3.2 Level 1 discovery

Level 1 discovery provides more thorough information about the TPer, and is performed using the `Properties control` session method. The reported properties mandatory for Opal are:

- `MaxMethods` — maximum number of methods per received subpacket. Default value is 1.
- `MaxSubpackets` — maximum number of subpackets per received packet. Default value is 1.
- `MaxPacketSize` — maximum size of a received packet. Default value is 1004.
- `MaxPackets` — maximum number of packets per received compacket. Default value is 1.
- `MaxComPacketSize` — maximum size of a received compacket. Default value is 1024.
- `MaxResponseComPacketSize` — maximum size of sent compacket.
- `MaxSessions` — maximum number of active sessions.
- `MaxIndTokenSize` — maximum size of an individual token. Default value is 968.
- `MaxAuthentications` — maximum possible authenticated authorities.

- MaxTransactionLimit —
- DefSessionTimeout, —

There are also defined properties which are optional for Opal:

- MaxReadSessions — maximum number of reader sessions.
- MaxAggTokenSize — maximum size of a combined token. Default value is 968.
- MaxSessionTimeout —
- MinSessionTimeout —
- DefTransTimeout —
- MaxTransTimeout —
- MinTransTimeout —
- MaxComIDTime — timeout for ComID.
- ContinuedTokens — . Default value is False.
- SequenceNumbers — . Default value is False.
- AckNak — . Default value is False.
- Asynchronous — . Default value is False.

Other than the defined properties, this discovery process may also report other, vendor specific, properties.

9.3.3 Level 2 discovery

Level 2 discovery is the act of reading any table and is provided by the Get method of an SP. This includes reading any table, such as the access control table or table of locking ranges. Some of the SPs' tables are described in later chapter ??.

9.4 Life cycle

In order to keep information about the state and the working capacity of an SP, the concept life cycle is introduced. Life cycle describes the condition of an SP using one of several states. In the Core standard the following states are introduced:

- “nonexistent” — the SP does not exist because it either was not created yet, or was already deleted.
- “issued” — the SP is in functional state.
- “issued-disabled” — the SP can only be authenticated to and enabled, all other functionality is disabled. An SP can be enabled and disabled using the SPInfo table of the corresponding SP.
- “issued-frozen” — no functionality of the SP is enabled. An SP can be frozen and unfrozen using the SP table of the Admin SP.
- “issued-disabled-frozen” — the SP is both disabled and frozen as described in the two previous points.
- “issued-failed” — SP is in fatal failure from which it cannot recover.

9.4.1 SP Issuance

SP issuance is the process of creation of a new SP from a base template and a set of other templates. After an SP is issued, it can be personalized. Personalization of an SP is the process of filling in tables, setting up authorities and access control.

9.4.2 Life cycles in Opal

Opal extends the states defined in the Core specification with a new set of life cycle states called “manufactured” — “manufactured-inactive”, “manufactured”, “manufactured-failed”, “manufactured-disabled”, “manufactured-disabled-frozen”, “manufactured-disabled”. Each of these new states mirrors the corresponding “issued” state from Core specification. However, compared to the “issued” states, the “manufactured” states are created during manufacturing by the manufacturer,

instead of during the subsequent use by the TPer owner. Because of this, the SPs are not issued and deleted in Opal, instead they are activated and reverted. Activation skips the process of personalization and the deletion maintains the existence of the SP. Note that this means that the “manufactured-inactive” corresponds to the “nonexistent” state.

9.5 Communication

In order to send commands to the TPer or the SPs and receive responses, a specific protocol must be used. The protocol is split into several layers. The high level method calls, optional transactions, and sessions, which are carried by the interface commands. ((It’s actually session, management, communication, TPer, interface, transport, in different part, fixup TODO.))

One of the information required in order to send a command is the ComID, a number identifying the caller (e.g. application of the host). It ensures that responses to method calls will be received by the correct application, since there can be at one time multiple callers. This number can be acquired through the level 0 discovery process.

The communication is based on packets. ComPacket is the primary one, each ComPacket contains data for communication under only a single ComID. Each ComPacket can contain several Packets. The usage of Packet is intended for session reliability, specifying sequential numbers and the acknowledgements for them. Every Packet then consists of one or more SubPackets. Finally, SubPacket contains one or more method calls, or results of method calls.

9.5.1 Method calls

The data of the SubPacket carries either the method call or method response. These are expressed using a set of tokens in a specific order.

structure is:

```
CALL_TOKEN | OBJECT_UID | METHOD_UID
| START_LIST_TOKEN | ARGUMENTS | END_LIST_TOKEN | END_OF_DATA_TOKEN
| START_LIST_TOKEN | ERROR_VALUE | END_LIST_TOKEN |
```

The general structure of response is identical to the structure of call, with different arguments used as a way of passing back information.

object uid can be SP, table, row of table of a special object uid
SMUID — session manager, used for session management

mandatory arguments:

optional arguments:

START_NAME_TOKEN | ARGUMENT_ID | ARGUMENT_VALUE | END_NAME_TOKEN

9.5.2 Sessions

Whereas ComID is used to differentiate the senders and therefore also the recipients of the methods, session is used for parallelization of communication of one such actor. Each session can have different authorized authorities and methods in process, but each session is still bound to one ComID. Since each ComID corresponds to one user, both the user and the TPer can send multiple methods calls or responses in one ComPacket.

TODO: not user or actor but host... e.g. if it would be network storage, there can be some thingie that would split the communication

Sessions are using system of readers-writer locks, to enable several concurrently running read sessions without causing concurrency issues.

There are two types of sessions: regular and control (not going to care about control much, just between TPer session manager and host session manager)... Each ComID has one associated control session with lifespan same as the ComID.

The control session is used for managing the regular sessions and as such provides methods such as:

- The `properties` method enables to find compromise between the communication capabilities of the host and the TPer, or to simply find the communication capabilities of the TPer. Using this method the properties such as maximum packet size are established.
- The `start session` and `sync session` methods provide a way to start an unencrypted session.

- The `start trusted session` and `sync trusted session` methods provide a way to start an encrypted session, after an unencrypted session has been established. The encrypted session provides secure messaging
- The `close session` method

*`trustedsession` - used with PuK, SymK, and HMAC authorities, secure messaging

TODO: talk about ComID lifecycle

9.5.3 Transactions

In order to facilitate safe execution of sequences of methods, the standard specifies transactions. Similarly to transactions in database systems, this feature enables one to revert effects of sequence of methods. This is done automatically in case the transaction is not finished, or if the transaction is manually aborted. However, not all the effects of methods are rolled back, such as logs. Nested transactions are supported, in which case the transaction is committed when the outermost transaction is finished.

9.5.4 Optional features

Note that some of the features described in the previous sections are optional. Even though every Opal device supports sessions, the minimum required maximum number of sessions is 1, ...

From our experience with Opal disks, most of them implemented only the required minimum <TODO: fact check later on>.

Note that even though in the previous sections it is said that there can be multiple methods per subpacket, or ... these features are all optional. Opal mandates only at least 1 sessions, 1 method per packet, 1 transaction, 1 packet per compacket, 1 subpacket per packet, and so on.

Even though the protocol specifies and can support many features such as secure messaging, asynchronous communication, or session reliability, these features are not mandatory for the Opal subsystem.

9.6 Features

move probably under templates,,, or even better SP?

9.6.1 Single User Mode

In many cases it might be desirable to prevent admins to access user data, even though the admins have generally more competence. In Opal this can be achieved using the Single User Mode [[tcg-sum](#)]. After activating Single User Mode, only a single User authority is capable of changing their authority object (this includes changing their PIN,), changing the LR proprieties (which includes the lock state of the LR), generate new key for the LR. The only action available to the admins are destructive actions.

9.6.2 MBR shadowing

This feature enables the disk to provide a fake master boot record (MBR). Instead of the MBR saved on the disk, the disk instead provides the saved shadow MBR on bootup. The shadow MBR may contain a software to authenticate the user to the disk and unlock it. After the shadow MBR is used, it can be deactivated by setting MBR done field in the MBRControl table.

The minimal maximum size of the shadow MBR seems to be 0x08000000 bytes (~134 MB).

Move to locking SP.

9.6.3 Access Control

In order to provide access to methods only to authorized actors, the standard also defines access control. The access control provides a way to allow access to methods of the SP only after a knowledge of a secret has been proved. Information required to verify knowledge of the secret is called a credential and is stored in one of the corresponding SP's credential table. Each SP may contain several credential tables, one for each type of credential, such as pin, RSA key pair, or AES key.

In order to support authentication of more than one user or require multiple users at once access control rules are defined using Access

Control Lists (ACL) containing Access Control Elements (ACE). Each entry in an ACL has an owner ACEs are Boolean expressions with inputs being authentication of authorities.

Other than an explicit authentication, an implicit authentication is also possible. In this authentication, the knowledge of secret is shown by successfully using encrypted communication channel. FACT CHECK

Each authority is an object in the authority table of the corresponding SP. Each authority has a type which is either individual or class. Class authorities correspond to a set of authorities, so that they can be easily assigned and changed in bulk.

Other than the Admin and User authorities there is also a special SID authority. This authority represents the owner of the TPer, and as such provides access of admins extended by the capabilities related to TPer management such as the ability to activate an SP or enabling Admin authorities. Compared to the other authorities, SID authority is shared between all SPs.

9.7 Templates

Since SPs may share some of their functionality such as authentication, modification of tables, or retrieval of data from tables, there exists templates to define this shared functionality. Each of the SP then implements a subset of functionality of one or more templates.

9.7.1 Base template

The base template defines the shared subset required by every SP, and is therefore mandatory. The most important functionality that is defined by this template is access control.

9.7.2 Admin template

Template specific for the unique Admin SP. Provides access to methods managing the TPer.

9.7.3 Crypto template

Template providing methods providing cryptographic methods.

9.7.4 Locking template

Template specific for the unique Locking SP. Provides access to methods managing the lock state of the device, locking ranges...

9.7.5 Clock template

Template providing methods for indicating current time, measuring lag, and access to monotonic counter.

9.7.6 Log template

Template providing methods for logging activity. The logging can be either carried out manually by an user (or an user application) or automatically by the TPer (as a result of invocations of methods of SPs containing this template, including invocations during read-only sessions). Only one system log table may exist on each SP, but multiple user log tables may be present on one SP. The logs are saved on non-volatile storage, depending on the security setting of the table, they may be either buffered or be saved after each entry.

TODO: how much are the logs preserved... will they get deleted with an SP?

9.8 Admin SP

holds information about the TPer and SPs, allows modification/creation/deletion of other SPs
unique

9.8.1 Methods

9.9 Locking SP

Locking SP procures the disk encryption and the locking and unlocking associated with it. This means that it provides access to manipulation of locking ranges, key generation and ...
unique

9.9.1 Locking range

The locking range feature gives the user a way to specify an LBA range on the disk that can be locked independently on the rest of the disk. Each locking range has also its own ACE, to control who can lock and unlock the range, and its own DEK.

There also always exists a special range called global locking range that is covering any area on disk that is not already covered by a regular locking range.

9.10 Implementations/Direct usage/something like this

9.10.1 Linux ioctl

Since version 4.11 Linux offers an ioctl requests to facilitate control over an Opal disk. Although these ioctl requests offer a simple access to control of the disk, not every feature is implemented. Some of the limitations are the following:

- Offers access to only 1 admin authority and up to 9 user authorities.
- no "write only" lock
- only basic binary table read (e.g. no iterations, no structured tables)

The individual requests for Opal functionality and their function are:

- SAVE — adds a key for a locking range into a unlock list, so that it can be used after waking the disk up from suspend. To wake the disk up, exported symbol `opal_unlock_from_suspend` can

be used to unlock the disk with the saved data. Note that the data is saved only in the RAM, and so this cannot be used to wake the computer.

- LOCK_UNLOCK — locks or unlocks reading or writing for the selected locking range.
- TAKE_OWNERSHIP — changes the admin authority password from the default one to the selected one. TODO: write about initialisation of the TPer to the generic chapter, C_PIN_MSID etc.
- ACTIVATE_LSP — changes state from "Manufactured-Inactive" to "Manufactured" state. Also facilitates setup of single user mode. TODO: write about TPer states in the generic chapter.
- SET_PW — changes the password of the selected authority, using the admin password.
- ACTIVATE_USR — enables specified user in the Opal tables.
- REVERT_TPR — reverts the TPer to the manufactured state using admin password.
- LR_SETUP — sets locking range position/locking enable
- ADD_USR_TO_LR — adds a user to the LR ACE. ...l_state == OPAL_RW means writing ace... (TODO: they seem to set it to "(user_uid || user_uid)"??? why? bug? -> MRE is opal_util_linux.c, only used last assigned to LR can control it, does (not) work on all test disks
- ENABLE_DISABLE_MBR — changes Enable parameter of the MBRControl. TODO: write about MBR in a generic chapter.
- ERASE_LR — calls the erase method, 00 00 06 00 00 08 03, can't find it in opal and in core it's "reserved for SSC", ... it's in single user mode standard — not only destroys data, also removes pin and user authority (for SUM).
- SECURE_ERASE_LR — regenerate the data encryption key of a range to destroy the previous one.

- `PSID_REVERT_TPR` — resets the TPer to the manufactured state using PSID.
- `MBR_DONE` — changes Done parameter of the `MBRControl`.
- `WRITE_SHADOW_MBR` — writes data into the MBR table.
- `GENERIC_TABLE_RW` — reads a table or writes into a byte table. Neither object tables nor iteration are supported.

Currently, there is no documentation to be found for the ioctl requests for Opal functionality. The information in the previous list was acquired from our code analysis.

9.10.2 Direct communication

Alternative to the OPAL ioctl are disk controller ioctls. Depending on the disk protocol, a different way of passing the OPAL commands is required, such as using `SG_IO` ioctl and `sg_io_hdr_t` structure for SCSI disks or `NVME_IOCTL_ADMIN_CMD` ioctl and `nvme_admin_cmd` structure for NVMe disks. Using these structures, the OPAL commands described in chapter ?? can be sent to the TPer. Compared to the OPAL ioctl this has the advantage of not being limited only to a subset of features that the OPAL ioctl implements, and instead being able to use every OPAL feature the device offers. This is not limited only to ... (e.g. improve performance by using optional features such as concurrency, etc, described earlier) ...

Although this gives the user access to every feature of the OPAL disk, it also requires them not only to implement the command hand-over for each type of disk separately, but also to create the methods and parse the method results, both described in chapter ??, on their own.

TODO disk interface popsat[NVME]

The commands that are called through the disk ioctl are called IF-RECV/IF-SEND in TCG Storage standards [4]. These commands corresponds to several other commands in different interface specification. Some of them are the Security Receive/Send commands in NVMe [8], the TRUSTED RECEIVE/SEND commands in ATA [9], or the SECURITY PROTOCOL IN/OUT commands in SCSI [10].

10 Notes

Write/read on locked range returns "Input/output error".
tested disks:

- SanDisk SD7UB2Q5
- Samsung SSD 850
 - problems with user authentication

ATA drives need `libata.allow_tpm` set to 1,,, but why?
templates — base, locking,, -> SPs. Service providers of the trusted peripherals are issued based on templates.

Global locking range is a special locking range that controls any area of the disk that is not controlled by any other locking range.

Shadow MBR allows FDE disk pre-boot. Presents “fake” MBR that can unlock the disk, on successful unlocking hands the control over to the real MBR.

key types (Credential Table Group):

- C_PIN — password
- C_RSA_* — signing arbitrary input???, session startup
- C_EC_* — ec-mqv, ec-dh session startup
- C_AES_*
- C_HMAC_*

5.3.4.1.3

authority operations:

- None — does not authenticate
- Password — authenticated using a pin
- Sign — challenge and response
- Exchange
- SymK

- HMAC
- TPerSign
- TPerExchange

Sessions may use secure messaging: use additionally a start/sync-trustedsession methods for challenge-response/key exchange

CVE-2018-12037 — “An issue was discovered on Samsung 840 EVO and 850 EVO devices (only in "ATA high" mode, not vulnerable in "TCG" or "ATA max" mode), ...” — wtf is ata high/ata max mode???

CVE-2018-12038

s
e
p
a
r
a
t
o
r

All that follows is from the core standard [4].

secure messaging — integrity/authentication and/or confidentiality,,,,

“Data confidentiality and access control over TPer features and capabilities: ... The protection provided by this exclusive access extends to confidentiality of instructions and data in transit between the trusted host application (or a TPM it uses) and the TPer”

secure messaging

sessions are using readers-writers lock.

SPs are combinations of templates, must contain base template (). other templates admin, clock, crypto, locking, log

“All SPs incorporate at least a subset of the Base Template’s tables and methods.” -> is \emptyset enough???

issuance = new SP from templates, personalization = customization of newly created SP (but can happen anytime...), initial data, authorities,

SSC = ... “A TPer MAY have only some of the capabilities (tables, methods, access controls, etc.) defined in this Core Specification and MAY include additional capabilities through table definitions and/or methods. A Security Subsystem Class SHALL NOT replace a capability called out in the Core Specification with the same capability implemented in different tables, methods, and access controls.”

stream encoding - methods and tokens...,TLV...

tables – bytes tables – raw data, just bytes (rows addressed by row number); object tables (rows addressed by uid, SP-unique, non-reusable, for anti-spoofing)

object = row of table

SP templates – base, admin, clock, crypto, locking, log,

interface – protocol-independent, IF-SEND, IF-RECV commands
COMID

ComID identifies the caller, each application has different ComID and therefore can communicate simultaneously, dynamic, static. Not a session, multiple sessions may use single comid. Then there are some states and transitions, active, associated, issued... Extended ComID for reusing ComID so that we can see if it is the old ComID or the new one (using extra two bytes).

PROTOCOL LAYERS

protocol layers — session, management, communication, tper, interface, transport (... the standard seems to imply that the comid is always managed??)

DISCOVERY

discovery levels : 3.3.5 : level 0 , level 1 Properties method of TPer, level 2 of SP

SESSIONS

sessions: regular and control (not going to care about control much, just between TPer session manager and host session manager); read/read-write mutexes...,

session manager methods - properties, start/sync session, start/sync trusted session, close session

*trustedsession - used with PuK, SymK, and HMAC authorities, secure messaging

authorities – used during session startup,,, host exchange authority – exchange of session keys, implicit authentication;;; host signing authority – challenge response authentication/startup method integrity, C_PIN (password), , , , same for host->SP

METHODS

TRANSACTIONS

— like in database, committing , etc.... optional (or implicit transaction on method level...), nested transaction committed when outermost finished, possible exceptions to rollbacks e.g. logs

Stream Flow Control – interface (handles sending IF-SEND/RECV commands across the interface) and stream data (handles not overwhelming host and TPer) types;;; uses Credit Control Subpacket to signify the opposite device that it is ready to receive data (and how much)

Session Reliability — ACKs, NACKs (SeqNumbers), timeouts,,, all optional

Synchronous Interface Communications — alternative to the previous asynchronous communication... to make it more simple,,, so IF-SEND to make TPer do something, IF-RECV to get the result, repeat ad nauseam

SP OPERATIONS

Special SP - admin SP – maintains info about TPer and other SPs

SP – Cartesian product of template (defining tables and methods, pretty sure I wrote about it earlier, but it is mentioned in the standard again) subsets,

ACCESS CONTROL

invoker may have to know secret (secret + public part = Credentials, operation of proving knowledge = Authentication Operation, proving knowledge = Authentication)

explicit authentication - e.g. password validation, challenge response,,, implicit authentication - e.g session key exchange

s

e

p

a

r

a

t

o

r

... And that's the introduction... what follows now is the method and table overview...

.....

10.0.1 Templates

The standard specifies several templates that can be used to create SPs.

Base Template –

Crypto Template — operates on Credential tables (the previously mentioned C_*, tables of other SPs,,,), can do enc, dec, sig, ver, hash, hmac, xor

Locking Template — provides access control

11 Data

11.0.1 Discovery

In order to collect information about Opal capabilities of the disks, we have written an utility program. This program uses the direct communication described in section ?? to communicate with the disk. We have decided to use the direct communication instead of the primarily because the capability to perform the discovery process using the Linux Opal ioctl is currently not possible. Even though there are patches being suggested, even if they were to be accepted, only the newest version of Linux kernel would support this feature, limiting the sample size. However, because this approach requires separate implementations of the disk interface commands, something about how this limits us only to a few disk interfaces because of testing, but it's not such a problem because they have a majority share, so find some source that this is actually true. The utility program performs the level 0 and level 1 discovery and the identify command to gather information about the disk. The aggregated and formatted output can be seen in table 11.1.

The first part of the table contains values acquired through the level 0 discovery, the second part (starting with MaxComPacketSize) contains values obtained through the level 2 discovery. In case the disk did not report a value, empty cell is used. There are some noticeable differences: since the first part is reported through firmly established C headers with static form, the numbers are parsed directly as C integers... the second part uses the TCG Storage protocol,

Table 11.1: Discovery process results, TODO: actually figure out how to do this nicely and easily, lol

[illegible]

with tokens that can be of different sizes — notably in `/dev/sdb` which returns numbers encoded with 4 bytes, even when not necessary. Some variables are dependent: because `MaxPackets` is always the minimum (1), then e.g `MaxPacketSize = MaxComPacketSize - (fixed size of MaxComPacketSize header)`.

Bibliography

1. *TCG Storage Security Subsystem Class: Opal*. 2022-01. Version 2.02. Standard. Trusted Computing Group. Available also from: <https://trustedcomputinggroup.org/resource/storage-work-group-storage-security-subsystem-class-opal>.
2. BIGGERS, Eric; TANGIRALA, Satya. *Inline Encryption*. 2021. Available also from: <https://kernel.org/doc/html/v5.17/block/inline-encryption.html>.
3. MÜLLER, Tilo; LATZO, Tobias; FREILING, Felix C. *Self-Encrypting Disks pose Self-Decrypting Risks: How to break Hardware-based Full Disk Encryption*. Available also from: <https://fau1-files.cs.fau.de/filepool/projects/sed/seds-at-risks.pdf>.
4. *TCG Storage Architecture Core Specification*. Tech. rep. Available also from: <https://trustedcomputinggroup.org/resource/tcg-storage-architecture-core-specification>.
5. *TCG Storage Security Subsystem Class: Opalite Specification*. Tech. rep. Available also from: <https://trustedcomputinggroup.org/resource/tcg-storage-security-subsystem-class-opalite>.
6. *TCG Storage Security Subsystem Class: Pyrite Specification*. Tech. rep. Available also from: <https://trustedcomputinggroup.org/resource/tcg-storage-security-subsystem-class-pyrite>.
7. *TCG Storage Security Subsystem Class: Ruby Specification*. Tech. rep. Available also from: <https://trustedcomputinggroup.org/resource/tcg-storage-security-subsystem-class-ruby-specification/>.
8. *NVM Express® Base Specification*. Tech. rep. Available also from: <https://nvmexpress.org/developers/nvme-specification/>.
9. *ATA/ATAPI Command Set - 3*. Tech. rep. Available also from: https://www.t10.org/ATAATAPI_Command_Set_-_3.pdf.
10. *SCSI Primary Commands - 4*. Tech. rep. Available also from: https://www.t10.org/members/w_spc4.htm.